

# Deep Learning: Homework1

1706 Wanru Zhao(17373240)

April 2020

## 1 Finding alternatives of softmax

$$\begin{aligned}\text{softmax}(x) &= \frac{e^x}{\sum e^x} \\ \text{Which of below work as alternative to Softmax?} \\ \text{abs-max}(x) &= \frac{|x|}{\sum |x|} \\ \text{square-max}(x) &= \frac{x^2}{\sum x^2} \\ \text{plus-one-abs-max}(x) &= \frac{1+|x|}{\sum 1+|x|} \\ \text{non-negative-max}(x) &= \frac{\max(0,x)}{\sum \max(0,x)}\end{aligned}$$

Figure 1: All Functions

### 1.1 Method

Since the `softmax_cross_entropy` function of tensorflow contains softmax inside, we cannot use the loss function of the baseline and need to modify the loss function according to

$$H(p, q) = - \sum_x p(x) \log q(x) \quad (1)$$

The code is as follows

```
1 loss = -tf.reduce_mean(  
2     tf.reduce_sum(tf.cast(label_onehot, dtype=tf.float32) *  
3     tf.log(tf.clip_by_value(preds, 1e-10, 1.0)), axis=-1))  
4 + loss_reg
```

Here, the `tf.clip_by_value()` method is used to smooth the predicted probability to ensure that there is no backward propagation of log operation with probability 0 in the result.

### 1.2 Results

Question	Function	Best Accuracy
baseline	soft-max	96.8%
q1.1	abs-max	28.7%
q1.2	square-max	43.8%
q1.3	plus-one-abs-max	95.1%
q1.4	non-negative-max	60.1%

accuracy




Figure 2: Baseline

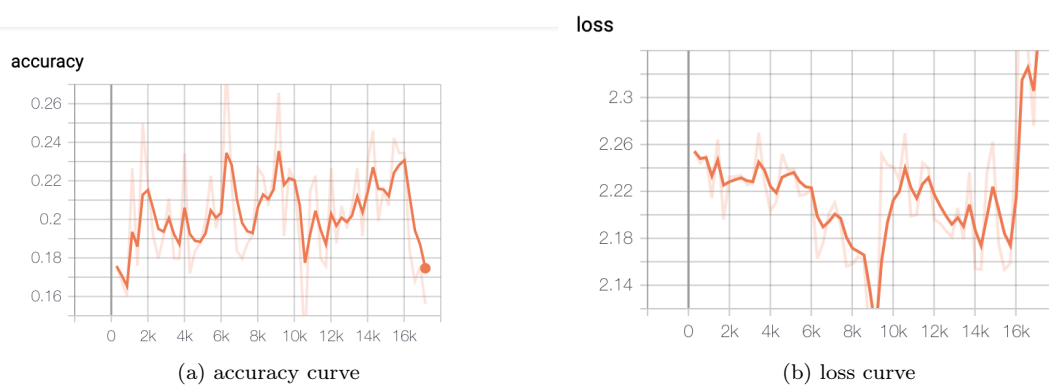


Figure 3: abs-max

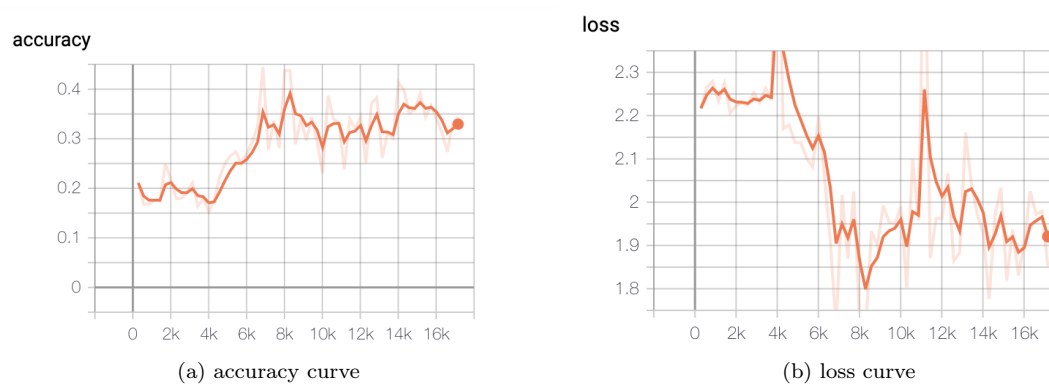


Figure 4: square-max




Figure 5: plus-one-abs-max

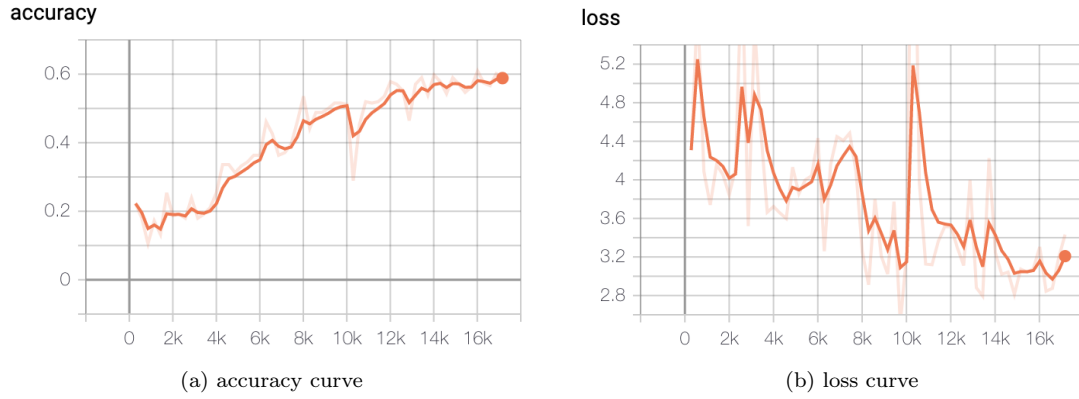


Figure 6: non-negative-max

### 1.3 Analysis

- 1 The softmax method of baseline enters into the cross entropy. The loss function is a convex function, and the greater the loss is, the greater the gradient is, which is convenient for back propagation and rapid convergence.
- 2 The abs - max method and square - max are even functions and not convex functions, so it's hard for them to converge.
- 3 The plus-one-abs-max method is the only one of the four methods that is close to the baseline. It can be an alternative of softmax.
- 4 The curve of the non-negative-max method is similar to that of the baseline, and it converges faster than the previous two methods, but less than the plus-one-abs-max method. It may also be an alternative of softmax.

## 2 Regression vs Classification

Change cross entropy loss to the square of euclidean distance between model predicted probability and one hot vector of the true label.

### 2.1 Method

Change the loss function to MSE according to

$$MSE(y, y') = \frac{\sum_{i=1}^n (y_i - y'_i)^2}{n} \quad (2)$$

```

1     if args.loss == 'regression':
2         loss = tf.reduce_mean(tf.reduce_sum(tf.square(preds -
3             tf.cast(label_onehot, dtype=tf.float32)), axis=1))
4     else:
5         loss = tf.losses.softmax_cross_entropy(label_onehot, logits) +
            loss_reg

```

## 2.2 Results

Question	Function	Accuracy
baseline	Cross Entropy Error	96.8%
q2	Mean Squared Error	95.1%

## 2.3 Analysis

The cross entropy loss is faster and converges better than the variance regression.

# 3 Lp pooling

Change all pooling layers to Lp pooling. [1]

## 3.1 Method

$$O = \left( \sum \sum I(i, j)^P * G(x, y) \right)^{\frac{1}{P}} \quad (3)$$

Lp pooling is a biologically inspired pooling layer modelled on complex cells who's operation can be summarized in equation (1), where G is a Gaussian kernel, I is the input feature map and O is the output feature map. It can be imagined as giving an increased weight to stronger features and suppressing weaker features. Two special cases of Lp pooling are notable. P = 1 corresponds to a simple Gaussian averaging, whereas P = ∞ corresponds to max-pooling (i.e only the strongest signal is activated).




Figure 7: a simple example of L2-pooling

Compute the gaussian kernel in advance.

```

1 def gauss(ksize):
2     sigma = ((ksize - 1) * 0.5 - 1) * 0.3 + 0.8
3     sum_val = 0
4     kernel = np.zeros((ksize, ksize))
5     center = ksize // 2
6     for i in range(ksize):
7         for j in range(ksize):
8             kernel[i, j] = np.exp(-((i - center) ** 2 + (j - center) **
9                 2) / (2 * sigma ** 2))
10            sum_val += kernel[i, j]
11     kernel = kernel / sum_val
12     return kernel

```

```

12 |
13 | G(x, y):
14 | [[0.05711826  0.12475775  0.05711826]
15 |  [0.12475775  0.27249597  0.12475775]
16 |  [0.05711826  0.12475775  0.05711826]]

```

The input feature dimension [WHC] is disassembled into  $c$  [wh1] features, and the corresponding [KKC] convolution is disassembled into  $c$  [kk1] gaussian kernel, each of which is put into the convolution function to generate  $c$  [wh1] output features, and spliced into [WHC] size. So we're convolving the characteristics of each channel. Finally, the stride length in convolution is set to 2, and finally the pooling effect is achieved by convolution.

### 3.2 Results

Question	Function	Accuracy
baseline	max-pooling	96.8%
q3	lp-pooling p=-1	92.1%
q3	lp-pooling p=1	93.4%
q3	lp-pooling p=2	96.2%
q3	lp-pooling p=4	95.1%

It can be seen from the experimental results that the lp-pooling test accuracy is significantly higher than that of the max-pooling at the baseline, indicating that the gaussian kernel selects better features than the maximum value after the convolution smoothing operation of the input features.

### 3.3 Analysis

## 4 Regularization

- Try Lp regularization with different p. Pick one number p with best accuracy.
- Set Lp regularization to a minus number. ( $L_{\text{model}} + L_{\text{reg}}$  to  $L_{\text{model}} - L_{\text{reg}}$ )

### 4.1 Method

When weight\_decay was equal to 1e-10, changing the P value made little difference to the result. So we set weight\_decay 1e-3.

```

1 | self.reg = lambda x: tf.reduce_sum(tf.pow(tf.abs(x), config.lp_reg))
2 |     * config.weight_decay

```

Set p=2 in q.2.

```

1 | self.reg = lambda x: tf.reduce_sum(tf.pow(tf.abs(x), config.lp_reg))
2 |     * config.weight_decay * (-1)

```

### 4.2 Results

Question	Method	Accuracy
q4.1	p=-1	82.8%
q4.1	p=1	90.9%
q4.1	p=2	93.6%
q4.1	p=4	95.4%

Method	Accuracy
baseline(no extra <sub>3232</sub> )	94.3%
q4.2(no extra <sub>3232</sub> )	94.3%

### 4.3 Analysis

- When  $p$  is negative, the results are worse. When  $p=4$ , the test accuracy is the highest. So the  $L_p4$  regularization constraint on the result is the most effective in this experiment.
- Whether the regularization is positive or negative has little effect on the accuracy of the whole neural network.

## References

- [1] Pierre Sermanet, Soumith Chintala, and Yann LeCun. Convolutional neural networks applied to house numbers digit classification. *CoRR*, abs/1204.3968, 2012.